

```

CREATE TABLE School (
    Name          VARCHAR(256) NOT NULL,
    TotalEnrollment  INTEGER,
    State          VARCHAR(20),
    StateSalaryRank  INTEGER,
    EarlyCareerPay   REAL,
    MidCareerPay     REAL,
    STEMPercent      INTEGER,
    Type            VARCHAR(20),
    DegreeLength     VARCHAR(10),
    BoardCost        REAL,
    InStateTuition   REAL,
    OutStateTuition  REAL
    PRIMARY KEY(Name)
);

```

```

CREATE TABLE Diversity (
    Name          VARCHAR(256) NOT NULL,
    WomenCount    INTEGER,
    AIANCount     INTEGER,
    AsianCount     INTEGER,
    BlackCount     INTEGER,
    HispanicCount  INTEGER,
    PacificCount   INTEGER,
    WhiteCount     INTEGER,
    PRIMARY KEY(Name),
    FOREIGN KEY(Name) REFERENCES School(Name) ON DELETE CASCADE
);

```

```

CREATE TABLE State (
    StateName      VARCHAR(20) NOT NULL,
    Population      INTEGER,
    CrimeRate       REAL,
    HappinessScore  REAL,
    Code            CHAR(2),
    PRIMARY KEY(StateName)
);

```

```

CREATE TABLE Company (
    CompanyName     VARCHAR(256) NOT NULL,
    Industry         VARCHAR(100),
    Location         VARCHAR(256),
    EmployeeCount    INTEGER,
    PRIMARY KEY(CompanyName)
);

```

);

CREATE TABLE RateProf (

StudentID	INTEGER NOT NULL,
ProfessorName	VARCHAR(50),
SchoolName	VARCHAR(256),
State	CHAR(2),
StarRating	REAL,
DifficultyRating	REAL,
CourseID	VARCHAR(20),
DepartmentName	VARCHAR(256),
PRIMARY KEY(StudentID)	

);

CREATE TABLE Major (

Major	VARCHAR(256) NOT NULL,
Ranking	INTEGER,
TotalCount	INTEGER,
WomenCount	INTEGER,
Category	VARCHAR(256),
EmploymentCount	INTEGER,
MedianSalary	REAL,
PRIMARY KEY(Major)	

);

Code version of DDL Commands:

```
CREATE TABLE School (Name VARCHAR(256) NOT NULL, TotalEnrollment INTEGER,
State VARCHAR(20), StateSalaryRank INTEGER, EarlyCareerPay REAL,
MidCareerPay REAL, STEMPercent INTEGER, Type VARCHAR(20), DegreeLength
VARCHAR(10), BoardCost REAL, InStateTuition REAL, OutStateTuition REAL,
PRIMARY KEY(Name));
```

```
CREATE TABLE Diversity (Name VARCHAR(256) NOT NULL, WomenCount INTEGER,
AIANCount INTEGER, AsianCount INTEGER, BlackCount INTEGER, HispanicCount
INTEGER, PacificCount INTEGER, WhiteCount INTEGER, PRIMARY KEY(Name),
FOREIGN KEY(Name) REFERENCES School(Name) ON DELETE CASCADE);
```

```
CREATE TABLE State (StateName VARCHAR(20) NOT NULL, Population INTEGER,
CrimeRate REAL, HappinessScore REAL, Code CHAR(2), PRIMARY KEY(StateName));
```

```
CREATE TABLE Company (CompanyName VARCHAR(256) NOT NULL, Industry
VARCHAR(100), Location VARCHAR(256), EmployeeCount INTEGER, PRIMARY
KEY(CompanyName));
```

```
CREATE TABLE RateProf (ProfessorName VARCHAR(50), SchoolName VARCHAR(256),
State CHAR(2), StarRating REAL, DifficultyRating REAL, CourseID
VARCHAR(20), DepartmentName VARCHAR(256), StudentID INTEGER NOT NULL,
PRIMARY KEY(StudentID));
```

```
CREATE TABLE Major (Ranking INTEGER, Major VARCHAR(256) NOT NULL,
TotalCount INTEGER, WomenCount INTEGER, Category VARCHAR(256),
EmploymentCount INTEGER, MedianSalary REAL, PRIMARY KEY(Major));
```

## Implemented Databases on GCP:

The screenshot displays the Google Cloud SQL console interface. The top navigation bar includes the Google Cloud logo, a dropdown menu for 'cs411-college-rec', a search bar, and various utility icons. The left sidebar shows the 'SQL' section with options for 'PRIMARY INSTANCE', 'Overview', 'Query insights', 'Connections', and 'Release Notes'. The main content area is titled 'Overview' and shows the instance 'college-rec' running MySQL 8.0. Below this, there is a chart for 'CPU utilization' over a 1-day period, showing a peak of 20%. At the bottom, a terminal window is open, displaying the following MySQL commands and output:

```
mysql> show tables;
ERROR 1046 (3D000): No database selected
mysql> use colleges;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> show tables;
+-----+
| Tables_in_colleges |
+-----+
| Company             |
| Diversity            |
| Major               |
| RateProf            |
| School              |
| State               |
+-----+
6 rows in set (0.01 sec)

mysql>
```

## 1000 rows of data in 3 tables

The image displays three screenshots of the Google Cloud SQL console, showing the 'Overview' page for a database instance named 'cs411-college-rec'. Each screenshot includes a CPU utilization chart and a terminal window with a SQL query and its results.

**Screenshot 1 (Top Left):** The terminal window shows the query `SELECT COUNT (Name) FROM School;` with the result `5536`.

**Screenshot 2 (Top Right):** The terminal window shows the query `SELECT COUNT (StudentID) FROM RateProf;` with the result `20000`.

**Screenshot 3 (Bottom):** The terminal window shows the query `SELECT COUNT (CompanyName) FROM Company;` with the result `355122`.

The CPU utilization charts in all three screenshots show a peak in utilization around 12:00 PM on Mar 8, reaching approximately 10%.

## Advanced SQL Query 1

This is a SQL query that finds the schools in states that have a **lower crime rate** than the average crime rate across the US. Order by CrimeRate in ascending order.

```
SELECT sc.Name, sc.State, st.CrimeRate FROM School sc JOIN State st ON  
(sc.State=st.Code) WHERE st.CrimeRate < (SELECT AVG(CrimeRate) FROM State) ORDER  
BY st.CrimeRate;
```

- Join
- Subqueries

The screenshot shows the Google Cloud SQL console. The top navigation bar includes the Google Cloud logo, a dropdown menu for 'cs411-college-rec', a search bar, and icons for help and settings. The main content area is divided into three sections: 'Overview', 'Need help connecting?', and 'New storage increase for instance'. The 'Overview' section shows the instance name 'cs411-college-rec' and various action buttons like 'EDIT', 'IMPORT', 'EXPORT', 'RESTART', 'STOP', 'DELETE', and 'CLONE'. The 'Need help connecting?' section provides links to documentation and a 'Learn more' link. The 'New storage increase for instance' section shows that automated backups and point-in-time recovery are enabled.

The terminal window at the bottom shows the following SQL query and its results:

```
mysql> SELECT sc.Name, sc.State, st.CrimeRate FROM School sc JOIN State st ON (sc.State=st.Code) WHERE st.CrimeRate < (SELECT AVG(CrimeRate) FROM State) ORDER BY st.CrimeRate LIMIT 15;
```

Name	State	CrimeRate
Antioch University New England	NH	2.344
Colby-Sawyer College	NH	2.344
College of St. Mary Magdalen	NH	2.344
Daniel Webster College	NH	2.344
Dartmouth College	NH	2.344
Franklin Pierce University	NH	2.344
Granite State College	NH	2.344
Great Bay Community College	NH	2.344
Keene State College	NH	2.344
Lakes Region Community College	NH	2.344
Manchester Community College (N.H.)	NH	2.344
N.H. Technical Institute-Concord's Community College	NH	2.344
Nashua Community College	NH	2.344
New England College	NH	2.344
New Hampshire Institute of Art	NH	2.344

15 rows in set (0.00 sec)

Before indexing, the cost is **9878.54**

```
mysql> EXPLAIN ANALYZE SELECT sc.Name, sc.State, st.CrimeRate FROM School sc JOIN State st ON (sc.State=st.Code) WHERE st.CrimeRate < (SELECT AVG(CrimeRate) FROM State) ORDER BY st.CrimeRate;
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Sort: st.CrimeRate (actual time=4.303..4.495 rows=2245 loops=1)
|   -> Stream results (cost=9878.54 rows=9852) (actual time=0.252..3.724 rows=2245 loops=1)
|     -> Inner hash join (sc.State = st.Code) (cost=9878.54 rows=9852) (actual time=0.246..2.996 rows=2245 loops=1)
|       -> Table scan on sc (cost=5.00 rows=5912) (actual time=0.048..1.785 rows=5536 loops=1)
|       -> Hash
|         -> Filter: (st.CrimeRate < (select #2)) (cost=1.92 rows=17) (actual time=0.146..0.160 rows=23 loops=1)
|           -> Table scan on st (cost=1.92 rows=50) (actual time=0.088..0.099 rows=50 loops=1)
|           -> Select #2 (subquery in condition; run only once)
|             -> Aggregate: avg(State.CrimeRate) (cost=10.25 rows=50) (actual time=0.040..0.040 rows=1 loops=1)
|               -> Table scan on State (cost=5.25 rows=50) (actual time=0.020..0.031 rows=50 loops=1)
|
+-----+
+-----+
```

## Design 1:

CREATE INDEX schoolstate\_idx on School(State);

The cost decreased from 9878.54 to **808.33** which is a big improvement! And it only took 1.47 seconds to access all the rows. This index seems important to keep.

```
mysql> EXPLAIN ANALYZE SELECT sc.Name, sc.State, st.CrimeRate FROM School sc JOIN State st ON (sc.State=st.Code) WHERE st.CrimeRate < (SELECT AVG(CrimeRate) FROM State) ORDER BY st.CrimeRate;
+-----+
| EXPLAIN |
+-----+
| -> Nested loop inner join (cost=808.33 rows=5685) (actual time=0.160..1.470 rows=2245 loops=1)
|   -> Sort: st.CrimeRate (cost=1.92 rows=50) (actual time=0.123..0.126 rows=23 loops=1)
|     -> Filter: ((st.CrimeRate < (select #2)) and (st.Code is not null)) (cost=1.92 rows=50) (actual time=0.094..0.104 rows=23 loops=1)
|       -> Table scan on st (cost=1.92 rows=50) (actual time=0.054..0.060 rows=50 loops=1)
|       -> Select #2 (subquery in condition; run only once)
|         -> Aggregate: avg(State.CrimeRate) (cost=10.25 rows=50) (actual time=0.024..0.024 rows=1 loops=1)
|           -> Table scan on State (cost=5.25 rows=50) (actual time=0.010..0.017 rows=50 loops=1)
|           -> Filter: (sc.State = st.Code) (cost=14.96 rows=114) (actual time=0.018..0.052 rows=98 loops=23)
|             -> Index lookup on sc using schoolstate_idx (State=st.Code) (cost=14.96 rows=114) (actual time=0.018..0.036 rows=98 loops=23)
|
| 1 row in set (0.01 sec)
```



```
CREATE INDEX state_idx on State(Code);
```

[illegible]

### Design 3:

```
CREATE INDEX crimerate_idx on State(CrimeRate);
```

After indexing, the cost is **902.11** for our query. This is worse. In fact, adding too many indexes can create overhead and slow down our performance. It seems to be more efficient to do a full table scan than create an index for CrimeRate, especially because we're only using CrimeRate to calculate the Avg(CrimeRate) and nothing else. It is not a very important/efficient index to keep.

```
mysql> EXPLAIN ANALYZE SELECT sc.Name, sc.State, st.CrimeRate FROM School sc JOIN State st ON (sc.State=st.Code) WHERE st.CrimeRate < (SELECT AVG(CrimeRate) FROM State) ORDER BY st.CrimeRate;
```

-----+  
| EXPLAIN  
  
-----+  
  
-----+  
|

-----+  
|-- Nested loop inner join (cost=902.11 rows=5685) (actual time=0.135..1.516 rows=2245 loops=1)  
--> Sort: st.CrimeRate (cost=5.25 rows=50) (actual time=0.059..0.062 rows=23 loops=1)  
-- Filter: ((st.CrimeRate < (select #2)) and (st.'Code' is not null)) (cost=5.25 rows=50) (actual time=0.021..0.032 rows=23 loops=1)  
--> Table scan on st (cost=5.25 rows=50) (actual time=0.016..0.023 rows=50 loops=1)  
-- Select #2 (subquery in condition; run only once)  
--> Aggregate: avg(State.CrimeRate) (cost=10.25 rows=50) (actual time=0.051..0.051 rows=1 loops=1)  
--> Index scan on State using crime\_rate\_idx (cost=5.25 rows=50) (actual time=0.035..0.044 rows=50 loops=1)  
-- Filter: (sc.State = st.'Code') (cost=14.77 rows=114) (actual time=0.023..0.056 rows=98 loops=23)  
--> Index lookup on sc using schoolstate\_idx (State=st.'Code') (cost=14.77 rows=114) (actual time=0.023..0.042 rows=98 loops=23)

In conclusion, we choose **Design 1** with the lowest cost and fastest lookup.

## Advanced SQL Query 2

This is a SQL query that finds the top schools with the highest average RateMyProfessor rating.

```
SELECT sc.Name, AVG(rp.StarRating) as avg_rating FROM School sc JOIN RateProf rp ON  
(sc.Name = rp.SchoolName) GROUP BY rp.SchoolName ORDER BY avg_rating DESC;
```

- Group By
- Join

```
mysql> SELECT sc.Name, AVG(rp.StarRating) as avg_rating FROM School sc JOIN RateProf rp ON (sc.Name = rp.SchoolName) GROUP BY rp.SchoolName ORDER BY avg_rating DESC LIMIT 15;
```

Name	avg_rating
Centralia College	5
Maryland Institute College of Art	5
Post University	5
Northwestern University	5
Clemson University	5
Springfield Technical Community College	5
Dakota Wesleyan University	5
Stonehill College	5
Averett University	4.9
Ohio State University: Newark Campus	4.8999999999999995
University of the Incarnate Word	4.8999999999999995
Bethel College	4.8000000000000001
Alma College	4.8
Washington University in St. Louis	4.8
Chabot College	4.8

```
15 rows in set (0.03 sec)
```

Before indexing, the cost is **8955.20**.

```
mysql> EXPLAIN ANALYZE SELECT sc.Name, AVG(rp.StarRating) as avg_rating FROM School sc JOIN RateProf rp ON (sc.Name = rp.SchoolName) GROUP BY rp.SchoolName ORDER BY avg_rating DESC;
+-----+
| EXPLAIN |
+-----+
|         |
+-----+
--> Sort: avg_rating DESC (actual time=33.110..33.145 rows=387 loops=1)
--> Table scan on <temporary> (actual time=0.002..0.119 rows=387 loops=1)
--> Aggregate using temporary table (actual time=32.823..32.965 rows=387 loops=1)
--> Nested loop inner join (cost=8955.20 rows=19811) (actual time=0.089..21.501 rows=13569 loops=1)
--> Filter: (rp.SchoolName is not null) (cost=2021.35 rows=9811) (actual time=0.054..8.004 rows=20000 loops=1)
--> Table scan on rp (cost=2021.35 rows=19811) (actual time=0.052..6.393 rows=20000 loops=1)
--> Single-row index lookup on sc using PRIMARY (Name=rp.SchoolName) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=20000)
|
+-----+
1 row in set (0.03 sec)
```

CREATE INDEX rating\_idx ON RateProf(StarRating);

The cost stayed the same at **8955.20**. This is not a very significant difference and the index may be unnecessary.

```
mysql> EXPLAIN ANALYZE SELECT sc.Name, AVG(rp.StarRating) as avg_rating FROM School sc JOIN RateProf rp ON (sc.Name = rp.SchoolName) GROUP BY rp.SchoolName ORDER BY avg_rating DESC;
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Sort: avg_rating DESC (actual time=32.184..32.215 rows=387 loops=1)
|   -> Table scan on <temporary> (actual time=0.002..0.096 rows=387 loops=1)
|     -> Aggregate using temporary table (actual time=31.916..32.033 rows=387 loops=1)
|       -> Nested loop inner join (cost=8955.20 rows=19811) (actual time=0.119..21.181 rows=13569 loops=1)
|         -> Filter: (rp.SchoolName is not null) (cost=2021.35 rows=19811) (actual time=0.079..7.929 rows=20000 loops=1)
|           -> Table scan on rp (cost=2021.35 rows=19811) (actual time=0.077..6.431 rows=20000 loops=1)
|             -> Single-row index lookup on sc using PRIMARY (Name=rp.SchoolName) (cost=0.25 rows=1) (actual time=0.000..0.001 rows=1 loops=20000)
|
+-----+
1 row in set (0.03 sec)
```

```
CREATE INDEX rating_idx ON RateProf(StarRating);
CREATE INDEX name_idx on School(Name);
```

The cost is now **8955.20** which shows that the cost remains the same. So the index appears to make no difference.

```
mysql> EXPLAIN ANALYZE SELECT sc.Name, AVG(rp.StarRating) as avg_rating FROM School sc JOIN RateProf rp ON (sc.Name = rp.SchoolName) GROUP BY rp.SchoolName ORDER BY avg_rating DESC;
+-----+
| EXPLAIN |
+-----+
|         |
+-----+
| -> Sort: avg_rating DESC (actual time=33.312..33.353 rows=387 loops=1)
|   -> Table Scan on <temporary> (actual time=0.003..0.106 rows=387 loops=1)
|     -> Aggregate using temporary table (actual time=33.035..33.163 rows=387 loops=1)
|       -> Nested loop inner join (cost=8955.20 rows=19811) (actual time=0.105..21.914 rows=13569 loops=1)
|         -> Filter: (rp.SchoolName is not null) (cost=2021.35 rows=19811) (actual time=0.065..8.189 rows=20000 loops=1)
|           -> Table scan on rp (cost=2021.35 rows=19811) (actual time=0.063..6.682 rows=20000 loops=1)
|             -> Single-row index lookup on sc using PRIMARY (Name=rp.SchoolName) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=20000)
|         |
|       |
|     |
|   |
+-----+
1 row in set (0.03 sec)
```

```
CREATE INDEX rating_idx ON RateProf(StarRating);
CREATE INDEX name_idx on School(Name);
CREATE INDEX rpschool_idx on RateProf(SchoolName);
```

The cost is now **8955.20** which shows that the cost remains the same. So the index appears to make no difference.

```
mysql> EXPLAIN ANALYZE SELECT sc.Name, AVG(rp.StarRating) as avg_rating FROM School sc JOIN RateProf rp ON (sc.Name = rp.SchoolName) GROUP BY rp.SchoolName ORDER BY avg_rating DESC;
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Sort: avg_rating DESC (actual time=33.866..33.900 rows=387 loops=1)
|   -> Table scan on <temporary> (actual time=0.003..0.116 rows=387 loops=1)
|     -> Aggregate using temporary table (actual time=33.576..33.715 rows=387 loops=1)
|       -> Nested loop inner join (cost=8955.20 rows=19811) (actual time=0.154..22.239 rows=13569 loops=1)
|         -> Filter: (rp.SchoolName is not null) (cost=2021.35 rows=19811) (actual time=0.087..8.344 rows=20000 loops=1)
|           -> Table scan on rp (cost=2021.35 rows=19811) (actual time=0.085..6.774 rows=20000 loops=1)
|             -> Single-row index lookup on sc using PRIMARY (Name=rp.SchoolName) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=20000)
|
+-----+
```

We conclude that no indexes are necessary. The default index is the simplest and efficient design.